

Credit Card Fraud Detection

Presented by Eduonix!

Throughout the financial sector, machine learning algorithms are being developed to detect fraudulent transactions. In this project, that is exactly what we are going to be doing as well. Using a dataset of nearly 28,500 credit card transactions and multiple unsupervised anomaly detection algorithms, we are going to identify transactions with a high probability of being credit card fraud. In this project, we will build and deploy the following two machine learning algorithms:

- Local Outlier Factor (LOF)
- Isolation Forest Algorithm

Furthermore, using metrics such as precision, recall, and F1-scores, we will investigate why the classification accuracy for these algorithms can be misleading.

In addition, we will explore the use of data visualization techniques common in data science, such as parameter histograms and correlation matrices, to gain a better understanding of the underlying distribution of data in our data set. Let's get started!

1. Importing Necessary Libraries

To start, let's print out the version numbers of all the libraries we will be using in this project. This serves two purposes - it ensures we have installed the libraries correctly and ensures that this tutorial will be reproducible.

In [1]:

```
import sys
import numpy
import pandas
import matplotlib
import seaborn
import scipy

print('Python: {}'.format(sys.version))
print('Numpy: {}'.format(numpy.__version__))
print('Pandas: {}'.format(pandas.__version__))
print('Matplotlib: {}'.format(matplotlib.__version__))
print('Seaborn: {}'.format(seaborn.__version__))
print('Scipy: {}'.format(scipy.__version__))
```

```
Python: 2.7.13 |Continuum Analytics, Inc.| (default, May 11 2017, 13:17:26) [MSC v.1500 64 bit (AMD64)]
Numpy: 1.14.0
Pandas: 0.21.0
Matplotlib: 2.1.0
Seaborn: 0.8.1
Scipy: 1.0.0
```

In [2]:

```
# import the necessary packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

2. The Data Set

In the following cells, we will import our dataset from a .csv file as a Pandas DataFrame. Furthermore, we will begin exploring the dataset to gain an understanding of the type, quantity, and distribution of data in our dataset. For this purpose, we will use Pandas' built-in describe feature, as well as parameter histograms and a correlation matrix.

In [3]:

```
# Load the dataset from the csv file using pandas
data = pd.read_csv('creditcard.csv')
```

In [4]:

```
# Start exploring the dataset
print(data.columns)
```

```
Index([u'Time', u'V1', u'V2', u'V3', u'V4', u'V5', u'V6', u'V7', u'V8', u'V9',
       u'V10', u'V11', u'V12', u'V13', u'V14', u'V15', u'V16', u'V17', u'V18',
       u'V19', u'V20', u'V21', u'V22', u'V23', u'V24', u'V25', u'V26', u'V27',
       u'V28', u'Amount', u'Class'],
      dtype='object')
```

In [5]:

```
# Print the shape of the data
data = data.sample(frac=0.1, random_state = 1)
print(data.shape)
print(data.describe())

# V1 - V28 are the results of a PCA Dimensionality reduction to protect user identities and sensitive features
```

```
(28481, 31)
```

	Time	V1	V2	V3	V4	\
count	28481.000000	28481.000000	28481.000000	28481.000000	28481.000000	
mean	94705.035216	-0.001143	-0.018290	0.000795	0.000350	
std	47584.727034	1.994661	1.709050	1.522313	1.420003	
min	0.000000	-40.470142	-63.344698	-31.813586	-5.266509	
25%	53924.000000	-0.908809	-0.610322	-0.892884	-0.847370	
50%	84551.000000	0.031139	0.051775	0.178943	-0.017692	
75%	139392.000000	1.320048	0.792685	1.035197	0.737312	
max	172784.000000	2.411499	17.418649	4.069865	16.715537	

	V5	V6	V7	V8	V9	\
count	28481.000000	28481.000000	28481.000000	28481.000000	28481.000000	
mean	-0.015666	0.003634	-0.008523	-0.003040	0.014536	
std	1.395552	1.334985	1.237249	1.204102	1.098006	
min	-42.147898	-19.996349	-22.291962	-33.785407	-8.739670	
25%	-0.703986	-0.765807	-0.562033	-0.208445	-0.632488	
50%	-0.068037	-0.269071	0.028378	0.024696	-0.037100	
75%	0.603574	0.398839	0.559428	0.326057	0.621093	
max	28.762671	22.529298	36.677268	19.587773	8.141560	

	V21	V22	V23	V24	\
count	28481.000000	28481.000000	28481.000000	28481.000000	
mean	0.004740	0.006719	-0.000494	-0.002626	
std	0.744743	0.728209	0.645945	0.603968	
min	-16.640785	-10.933144	-30.269720	-2.752263	
25%	-0.224842	-0.535877	-0.163047	-0.360582	
50%	-0.029075	0.014337	-0.012678	0.038383	
75%	0.189068	0.533936	0.148065	0.434851	
max	22.588989	6.090514	15.626067	3.944520	

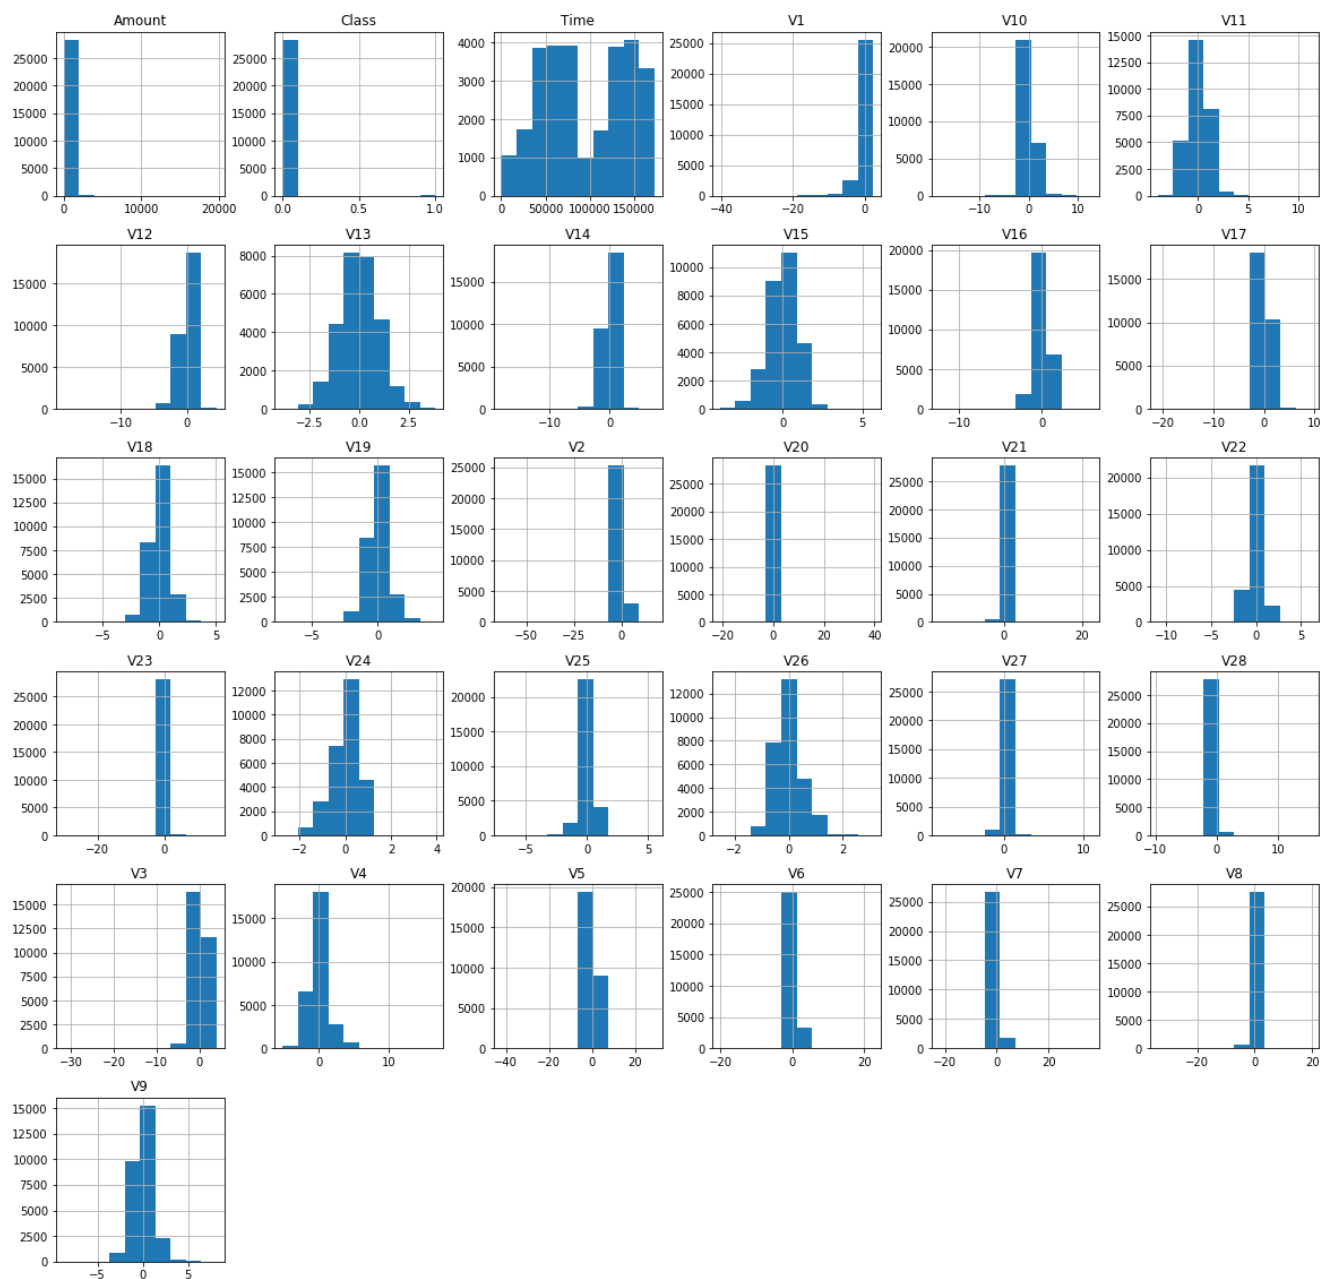
	V25	V26	V27	V28	Amount	\
count	28481.000000	28481.000000	28481.000000	28481.000000	28481.000000	
mean	-0.000917	0.004762	-0.001689	-0.004154	89.957884	
std	0.520679	0.488171	0.418304	0.321646	270.894630	
min	-7.025783	-2.534330	-8.260909	-9.617915	0.000000	
25%	-0.319611	-0.328476	-0.071712	-0.053379	5.980000	
50%	0.015231	-0.049750	0.000914	0.010753	22.350000	
75%	0.351466	0.253580	0.090329	0.076267	78.930000	
max	5.541598	3.118588	11.135740	15.373170	19656.530000	

	Class
count	28481.000000
mean	0.001720
std	0.041443
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 31 columns]

In [6]:

```
# Plot histograms of each parameter
data.hist(figsize = (20, 20))
plt.show()
```



In [7]:

```
# Determine number of fraud cases in dataset

Fraud = data[data['Class'] == 1]
Valid = data[data['Class'] == 0]

outlier_fraction = len(Fraud)/float(len(Valid))
print(outlier_fraction)


print('Fraud Cases: {}'.format(len(data[data['Class'] == 1])))
print('Valid Transactions: {}'.format(len(data[data['Class'] == 0])))
```

```
0.00172341024198
Fraud Cases: 49
Valid Transactions: 28432
```

In [8]:

```
# Correlation matrix
corrmat = data.corr()
fig = plt.figure(figsize = (12, 9))

sns.heatmap(corrmat, vmax = .8, square = True)
plt.show()
```



In [9]:

```
# Get all the columns from the dataframe
columns = data.columns.tolist()

# Filter the columns to remove data we do not want
columns = [c for c in columns if c not in ["Class"]]

# Store the variable we'll be predicting on
target = "Class"

X = data[columns]
Y = data[target]

# Print shapes
print(X.shape)
print(Y.shape)
```

```
(28481, 30)
(28481L,)
```

3. Unsupervised Outlier Detection

Now that we have processed our data, we can begin deploying our machine learning algorithms. We will use the following techniques:

Local Outlier Factor (LOF)

The anomaly score of each sample is called Local Outlier Factor. It measures the local deviation of density of a given sample with respect to its neighbors. It is local in that the anomaly score depends on how isolated the object is with respect to the surrounding neighborhood.

Isolation Forest Algorithm

The IsolationForest 'isolates' observations by randomly selecting a feature and then randomly selecting a split value between the maximum and minimum values of the selected feature.

Since recursive partitioning can be represented by a tree structure, the number of splittings required to isolate a sample is equivalent to the path length from the root node to the terminating node.

This path length, averaged over a forest of such random trees, is a measure of normality and our decision function.

Random partitioning produces noticeably shorter paths for anomalies. Hence, when a forest of random trees collectively produce shorter path lengths for particular samples, they are highly likely to be anomalies.

In [11]:

```
from sklearn.metrics import classification_report, accuracy_score
from sklearn.ensemble import IsolationForest
from sklearn.neighbors import LocalOutlierFactor

# define random states
state = 1

# define outlier detection tools to be compared
classifiers = {
    "Isolation Forest": IsolationForest(max_samples=len(X),
                                         contamination=outlier_fraction,
                                         random_state=state),
    "Local Outlier Factor": LocalOutlierFactor(
        n_neighbors=20,
        contamination=outlier_fraction)}
```

In [15]:

```
# Fit the model
plt.figure(figsize=(9, 7))
n_outliers = len(Fraud)

for i, (clf_name, clf) in enumerate(classifiers.items()):

    # fit the data and tag outliers
    if clf_name == "Local Outlier Factor":
        y_pred = clf.fit_predict(X)
        scores_pred = clf.negative_outlier_factor_
    else:
        clf.fit(X)
        scores_pred = clf.decision_function(X)
        y_pred = clf.predict(X)

    # Reshape the prediction values to 0 for valid, 1 for fraud.
    y_pred[y_pred == 1] = 0
    y_pred[y_pred == -1] = 1

    n_errors = (y_pred != Y).sum()

    # Run classification metrics
    print('{}: {}'.format(clf_name, n_errors))
    print(accuracy_score(Y, y_pred))
    print(classification_report(Y, y_pred))
```

Local Outlier Factor: 97

0.9965942207085425

	precision	recall	f1-score	support
0	1.00	1.00	1.00	28432
1	0.02	0.02	0.02	49
avg / total	1.00	1.00	1.00	28481

Isolation Forest: 71

0.99750711000316

	precision	recall	f1-score	support
0	1.00	1.00	1.00	28432
1	0.28	0.29	0.28	49

avg / total	1.00	1.00	1.00	28481
-------------	------	------	------	-------